

“Client-Server”; Enforced Modularity

Key Concepts:

1. The limitations of convention-based “soft” modularity
2. The high-level approach of enforced modularity
 - a. What are the challenges with implementing enforced modularity?
 - b. What benefits do we get from enforced modularity?
3. Micro-services—enforced modularity in the modern world
4. RPCs – the abstraction, the implementation, their benefits, their limitations

Motivation: The limits of “soft” modularity

- Modularity is insanely broad. The boxes and arrows might define:
 - *Soft*: Boxes are functions, classes, arrows are invocations
 - *Enforced/hard*: Boxes are computers, arrows are communication links
- More often than not, we practice soft modularity.
 - It might be ALL that you have practiced up to this point!
 - Why is this a problem???
- Concrete Example: Function call that gets a password:
- Review: how do function calls work?
 - The specifics of this process depend on the architecture's *calling conventions*:
 - Such conventions are enforced by... compilers, runtimes, etc.
 - Existing systems are variations on a theme: They all save return addresses on the stack and store local variables on stack, though.
 - What can go wrong?
 - What happens if gets receives a huge value?
 - Overwrite other local variables – Data Corruption
 - Overwrite return address – Return Oriented Programming
 - Claim: convention-based modularity is fundamentally brittle and insecure
- **The Crux of the problem: caller and callee share a memory**
 - Question: if we gave them separate naming contexts, could the same problems occur? Depends: can you access an object outside of your context?
 - Yields large propagation of effects:
 - Spurious writes (accidental or malicious)
 - Arbitrary access to modify global variables (**issue with local copies**)
- **The Crux of the problem: caller and callee share an interpreter:**
 - No control over a “runaway module” [prop of effects]:
 - What if get_input executes forever? (infinite loop)
 - No module-level permissions/security [prop of effects]:
 - Validate and get_input fundamentally share a protection level.
 - Cannot manage performance [incommensurate scaling]:
 - The performance of validate is tied to get_input.
 - Modules share a fate in the presence of failures:

- If `get_input` fails, so does `validate`.
- Impedes specificity in *fault tolerant* (or the idea that the system should survive certain classes of faults)

Could you solve these issues with better languages?

- High-level languages constrain memory sharing (e.g., java, Go, etc.)
- People have tried implementing real systems with these:
 - Biscuit [OSDI 2018] built an operating system entirely in Go
 - Singularity [Microsoft 2000s] built an OS in C#
- Problem: High-Level languages don't seem to solve the interpreter issues
- Problem: High Level languages require expensive runtime support
- Problem: Should you **really** trust these languages?
 - What if they have accidental bugs?
 - What if they are malicious actors?
 - Ken Thompson's Reflections on Trusting Trust
 - Such injections **are** possible! (see <https://www.theverge.com/2021/4/30/22410164/linux-kernel-university-of-minnesota-banned-open-source>)

Enforced/Hard Modularity

- What if we place them in entirely separate contexts?
- Separate contexts increase *isolation* of modules
 - Independent Memory: No way to modify another modules memory directly
 - Cannot clobber other module's memory
 - Cannot write to other module's instruction reference
 - Independent Interpreters:
 - Contain runaway module (assuming timeouts)
 - Per-interpreter permissions
 - Enable independent scaling
 - Enable less fate sharing and more interesting fault tolerance
 - But, what if they need to talk? Communication links! (Message Passing)
 - Synergistic benefit: Interchangeability. Think of all of the web clients!
- Challenges:
 - Representation: How do computers speak the same language?
 - E.g., little vs. big endian.
 - E.g., "self-contained" messages: messages need to include not just arguments, but also information about specific function, maybe even which client is called (e.g., HTTP put vs. HTTP get)
 - *marshaling* or *serialization*: The processes of converting internal binary data into an externally consumable format
 - You have to parse each message! No compiler :(
 - You *get* to handle failures. But also, you *must* consider them!

- Comparatively difficult to change a protocol
- Benefits:
 - Scalability
 - Better fault tolerance-performance tradeoffs
 - Trust

RPC: reducing the developer burden of using Client-Server

- Abstraction—execute some code on another computer
- Implementation:
 - REST:
 - Protobuf, Apache Thrift, etc.
- This seems great! Abstraction to the rescue, right!
 - Developers never have to serialize again!
 - All the peculiarities of message-passing are hidden from the user!
- Problem: is RPC the same semantics as functions? Why or why not?
 - Performance! Serialization is... Expensive. Unpredictable.
 - Semantics! Functions give you *exactly once semantics*. Can RPC?

Summary Table

Isolation	Soft Modularity easy + Conventions	Hard Modularity Physical Boundaries + MP
Mechanism	functions classes	clients + servers
failures	total	Partial.
efficiency	depends	depends. But possibly worse
Programmability	easy!	more difficult (passing. Faulted failures)