

# Deadlocks



# Deadlocks: overview

---

- ❖ Resources
- ❖ Why do deadlocks occur?
- ❖ Dealing with deadlocks
  - ▶ Ignoring them: ostrich algorithm
  - ▶ Detecting & recovering from deadlock
  - ▶ Avoiding deadlock
  - ▶ Preventing deadlock

# Resources

- ❖ Resource: something a process or thread uses
  - ▶ Usually limited (at least somewhat)
- ❖ Examples of computer resources
  - ▶ Printers
  - ▶ Semaphores / locks
  - ▶ Memory
  - ▶ Database tables
- ❖ Processes need access to resources in a reasonable order
- ❖ Two types of resources:
  - ▶ Preemptable resources: can be taken away from a process with no ill effects
  - ▶ Nonpreemptable resources: will cause the process to fail if taken away

# Using resources

- ❖ Sequence of events required to use a resource
  - ▶ Request the resource
  - ▶ Use the resource
  - ▶ Release the resource
- ❖ Can't use the resource if request is denied
  - ▶ Requesting process has options
    - Block and wait for resource
    - Continue (if possible) without it: may be able to use an alternate resource
    - Process fails with error code
  - ▶ Some of these may be able to prevent deadlock...


# When do deadlocks happen?

## ❖ Suppose

- ▶ Process 1 holds resource A and requests resource B
- ▶ Process 2 holds B and requests A
- ▶ Both can be blocked, with neither able to proceed

## ❖ Deadlocks occur when ...

- ▶ Processes are granted exclusive access to devices or software constructs (resources)
- ▶ Each deadlocked process needs a resource held by another deadlocked process



# What is a deadlock?

- ❖ Formal definition:
  - “A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”
- ❖ Usually, the event is release of a currently-held resource
- ❖ In deadlock, none of the processes can
  - ▶ Run
  - ▶ Release resources
  - ▶ Be awakened




# Four conditions for deadlock

---

- ❖ Mutual exclusion
  - ▶ Each resource is assigned to at most one process
- ❖ Hold and wait
  - ▶ A process holding resources can request more resources
- ❖ No preemption
  - ▶ Previously granted resources cannot be forcibly taken away
- ❖ Circular wait
  - ▶ There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain



# Resource allocation graphs



- ❖ Resource allocation modeled by directed graphs
- ❖ Example 1:
  - Resource R assigned to process A
- ❖ Example 2:
  - Process B is requesting / waiting for resource S
- ❖ Example 3:
  - Process C holds T, waiting for U
  - Process D holds U, waiting for T
  - C and D are in deadlock!





# Dealing with deadlock

- ❖ How can the OS deal with deadlock?
  - ▶ Ignore the problem altogether!
    - Hopefully, it'll never happen...
  - ▶ Detect deadlock & recover from it
  - ▶ Dynamically avoid deadlock
    - Careful resource allocation
  - ▶ Prevent deadlock
    - Remove at least one of the four necessary conditions
  
- ❖ We'll explore these tradeoffs

# Getting into deadlock



## A

Acquire R  
Acquire S  
Release R  
Release S





## B

Acquire S  
Acquire T  
Release S  
Release T



## C

Acquire T  
Acquire R  
Release T  
Release R



# The Ostrich Algorithm

---

- ❖ Pretend there's no problem
- ❖ Reasonable if
  - ▶ Deadlocks occur very rarely
  - ▶ Cost of prevention is high
- ❖ UNIX™ and Windows™ take this approach
  - ▶ Resources (memory, CPU, disk space) are plentiful
  - ▶ Deadlocks over such resources rarely occur
  - ▶ Deadlocks typically handled by rebooting
- ❖ Trade off between convenience and correctness

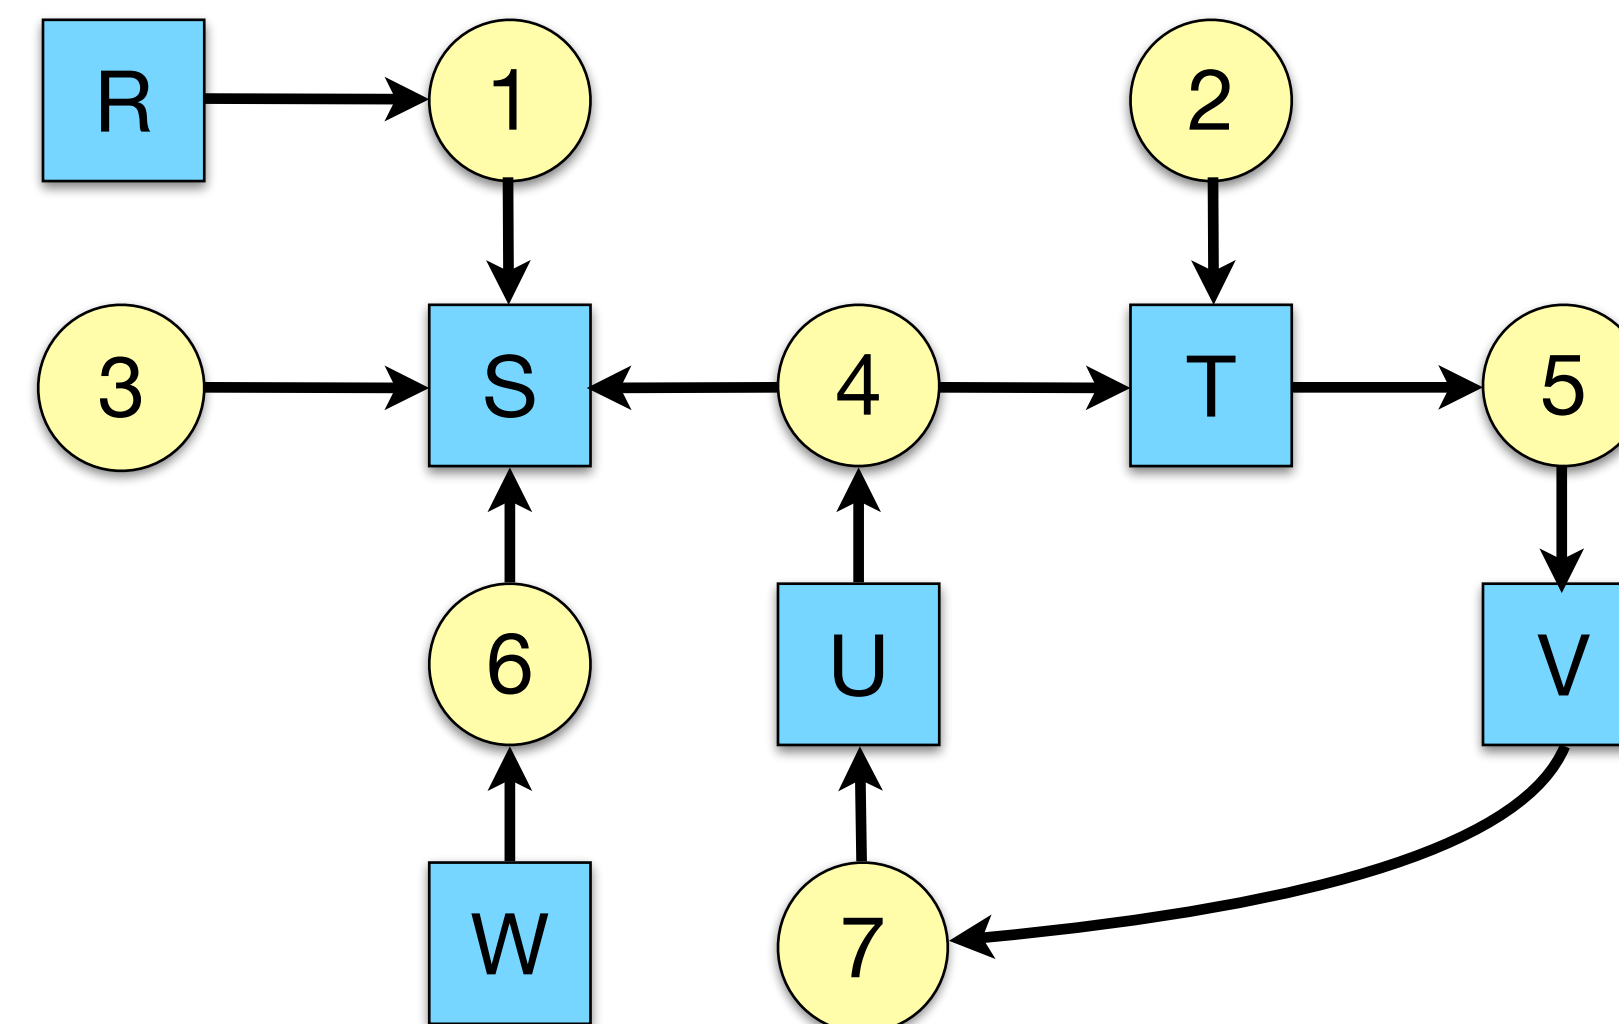
# Not getting into deadlock...

- ❖ Many situations may result in deadlock (but don't have to)
  - ▶ In previous example, A could release R before C requests R, resulting in no deadlock
  - ▶ Can we always get out of it this way?
- ❖ Find ways to:
  - ▶ Detect deadlock and reverse it
  - ▶ Stop it from happening in the first place

# Detecting deadlocks using graphs

- ❖ Process holdings and requests in the table and in the graph (they're equivalent)
- ❖ Graph contains a cycle  $\Rightarrow$  deadlock!
  - ▶ Easy to pick out by looking at it (in this case)
  - ▶ Need to mechanically detect deadlock
- ❖ Not all processes are deadlocked (1, 3, 6 not in deadlock)

Process	Holds	Wants
1	R	S
2		T
3		S
4	U	S,T
5	T	V
6	W	S
7	V	U



# Recovering from deadlock

- ❖ Recovery through preemption
  - ▶ Take a resource from some other process
  - ▶ Depends on nature of the resource and the process
- ❖ Recovery through rollback
  - ▶ Checkpoint a process periodically
  - ▶ Use saved state to restart the process if it's in deadlock
  - ▶ May present a problem if the process affects lots of “external” things
- ❖ Recovery through killing processes
  - ▶ Crudest but simplest way to break a deadlock: kill one of the processes in the deadlock cycle
  - ▶ Other processes can get its resources
  - ▶ Try to choose a process that can be rerun from the start
    - Pick one that hasn't run too far already

# Preventing deadlock

---

- ❖ Deadlock can be completely prevented!
- ❖ Ensure that at least one of the conditions for deadlock never occurs
  - ▶ Mutual exclusion
  - ▶ Circular wait
  - ▶ Hold & wait
  - ▶ No preemption
- ❖ Not always possible...



# Eliminating mutual exclusion

---


- ❖ Some devices (such as printer) can be spooled
  - ▶ Only the printer daemon uses printer resource
  - ▶ This eliminates deadlock for printer
- ❖ Not all devices can be spooled
- ❖ Principle:
  - ▶ Avoid assigning resource when not absolutely necessary
  - ▶ As few processes as possible actually claim the resource

# Attacking “hold and wait”

- ❖ Require processes to request resources before starting
  - ▶ A process never has to wait for what it needs
- ❖ This can present problems
  - ▶ A process may not know required resources at start of run
  - ▶ This also ties up resources other processes could be using
    - Processes will tend to be conservative and request resources they might need
- ❖ Variation: a process must give up all resources before making a new request
  - ▶ Process is then granted all prior resources as well as the new ones
  - ▶ Problem: what if someone grabs the resources in the meantime—how can the process save its state?


# Attacking “no preemption”

- ❖ This is not usually a viable option
- ❖ Consider a process given the printer
  - ▶ Halfway through its job, take away the printer
  - ▶ Confusion ensues!
- ❖ May work for some resources
  - ▶ Forcibly take away memory pages, suspending the process
  - ▶ Process may be able to resume with no ill effects



# Attacking “circular wait”

- ❖ Assign an order to resources
- ❖ Always acquire resources in numerical order
  - ▶ Need not acquire them all at once!
- ❖ Circular wait is prevented
  - ▶ A process holding resource  $n$  can't wait for resource  $m$  if  $m < n$
  - ▶ No way to complete a cycle!
    - Place processes above the highest resource they hold and below any they're requesting
    - All arrows point up!



# Deadlock prevention: summary

Condition	Prevented by
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away if there's not a complete set
Circular wait	Order resources numerically